

## Employee Attrition Prediction

Employee attrition prediction is a data-driven approach to identify factors that lead to employees leaving a company and forecasting which employees are at risk of attrition. By analyzing historical employee data, organizations can develop predictive models to improve retention strategies, reduce turnover costs, and enhance workplace satisfaction.

### Import the necessary libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, ConfusionMatrixDisplay, accuracy_score
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.ensemble import RandomForestRegressor
import pickle
import warnings
```

```
warnings.filterwarnings('ignore')
```


### Read the dataset

```
df=pd.read_csv('/content/drive/MyDrive/train_main.csv')
df
```



	Employee ID	Age	Gender	Years at Company	Job Role	Monthly Income	Work-Life Balance	Job Satisfaction	Performance Rating
0	8410	31	Male	19	Education	5390	Excellent	Medium	A
1	64756	59	Female	4	Media	5534	Poor	High	
2	30257	24	Female	10	Healthcare	8159	Good	High	
3	65791	36	Female	7	Education	3989	Good	High	
4	65026	56	Male	41	Education	4821	Fair	Very High	A
...	...	...	...	...	...	...	...	...	
59593	37195	50	Female	12	Education	4414	Fair	High	A
59594	6266	18	Male	4	Healthcare	8040	Fair	High	
59595	54887	22	Female	14	Technology	7944	Fair	High	
59596	861	23	Male	8	Education	2931	Fair	Very High	A
59597	15796	56	Male	19	Technology	6660	Good	High	A

59598 rows × 24 columns



df.columns



```
Index(['Employee ID', 'Age', 'Gender', 'Years at Company', 'Job Role',  
      'Monthly Income', 'Work-Life Balance', 'Job Satisfaction',  
      'Performance Rating', 'Number of Promotions', 'Overtime',  
      'Distance from Home', 'Education Level', 'Marital Status',  
      'Number of Dependents', 'Job Level', 'Company Size', 'Company Tenure',  
      'Remote Work', 'Leadership Opportunities', 'Innovation Opportunities',  
      'Company Reputation', 'Employee Recognition', 'Attrition'],  
      dtype='object')
```

Check for the missing values

```
df.isna().sum()
```



	0
Employee ID	0
Age	0
Gender	0
Years at Company	0
Job Role	0
Monthly Income	0
Work-Life Balance	0
Job Satisfaction	0
Performance Rating	0
Number of Promotions	0
Overtime	0
Distance from Home	0
Education Level	0
Marital Status	0
Number of Dependents	0
Job Level	0
Company Size	0
Company Tenure	0
Remote Work	0
Leadership Opportunities	0
Innovation Opportunities	0
Company Reputation	0
Employee Recognition	0
Attrition	0

dtype: int64

df.dtypes



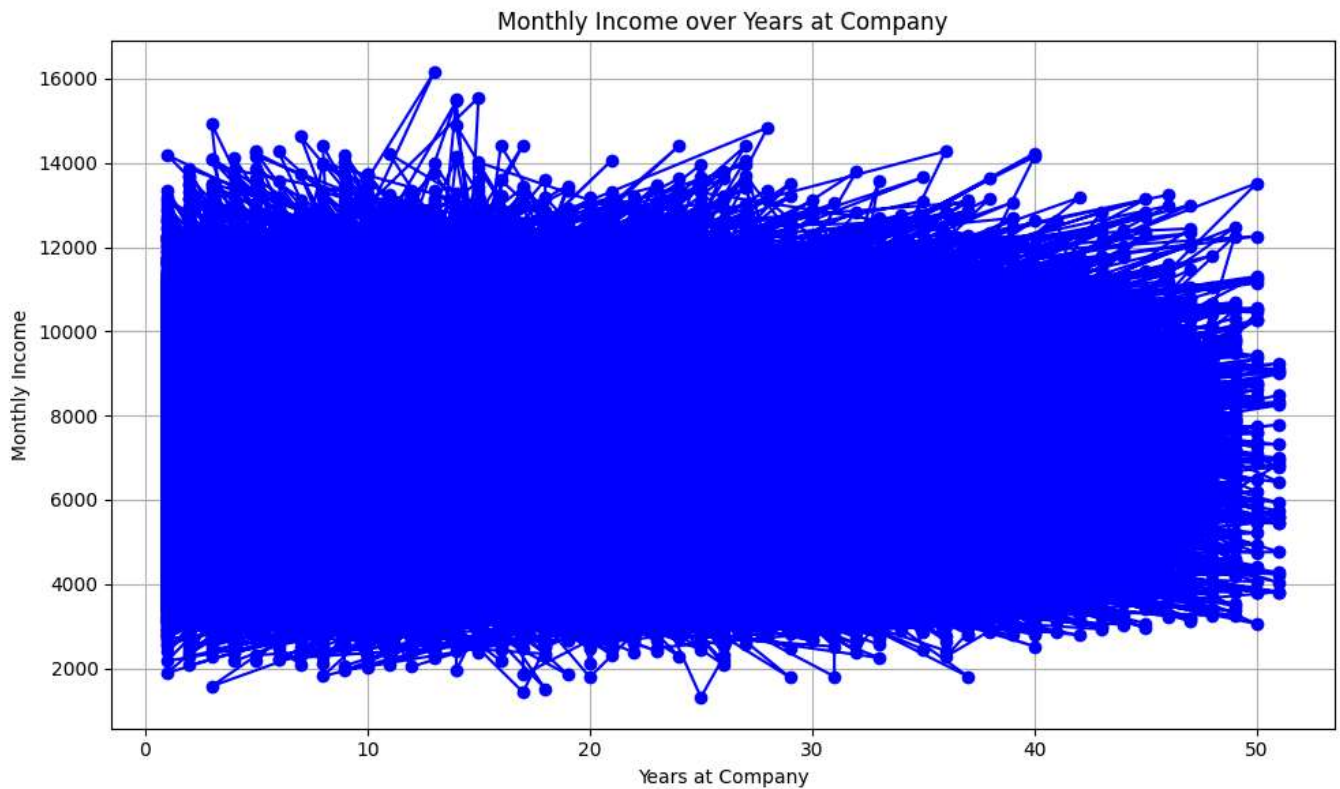
0

<b>Employee ID</b>	int64
<b>Age</b>	int64
<b>Gender</b>	object
<b>Years at Company</b>	int64
<b>Job Role</b>	object
<b>Monthly Income</b>	int64
<b>Work-Life Balance</b>	object
<b>Job Satisfaction</b>	object
<b>Performance Rating</b>	object
<b>Number of Promotions</b>	int64
<b>Overtime</b>	object
<b>Distance from Home</b>	int64
<b>Education Level</b>	object
<b>Marital Status</b>	object
<b>Number of Dependents</b>	int64
<b>Job Level</b>	object
<b>Company Size</b>	object
<b>Company Tenure</b>	int64
<b>Remote Work</b>	object
<b>Leadership Opportunities</b>	object
<b>Innovation Opportunities</b>	object
<b>Company Reputation</b>	object
<b>Employee Recognition</b>	object
<b>Attrition</b>	object

**dtype:** object

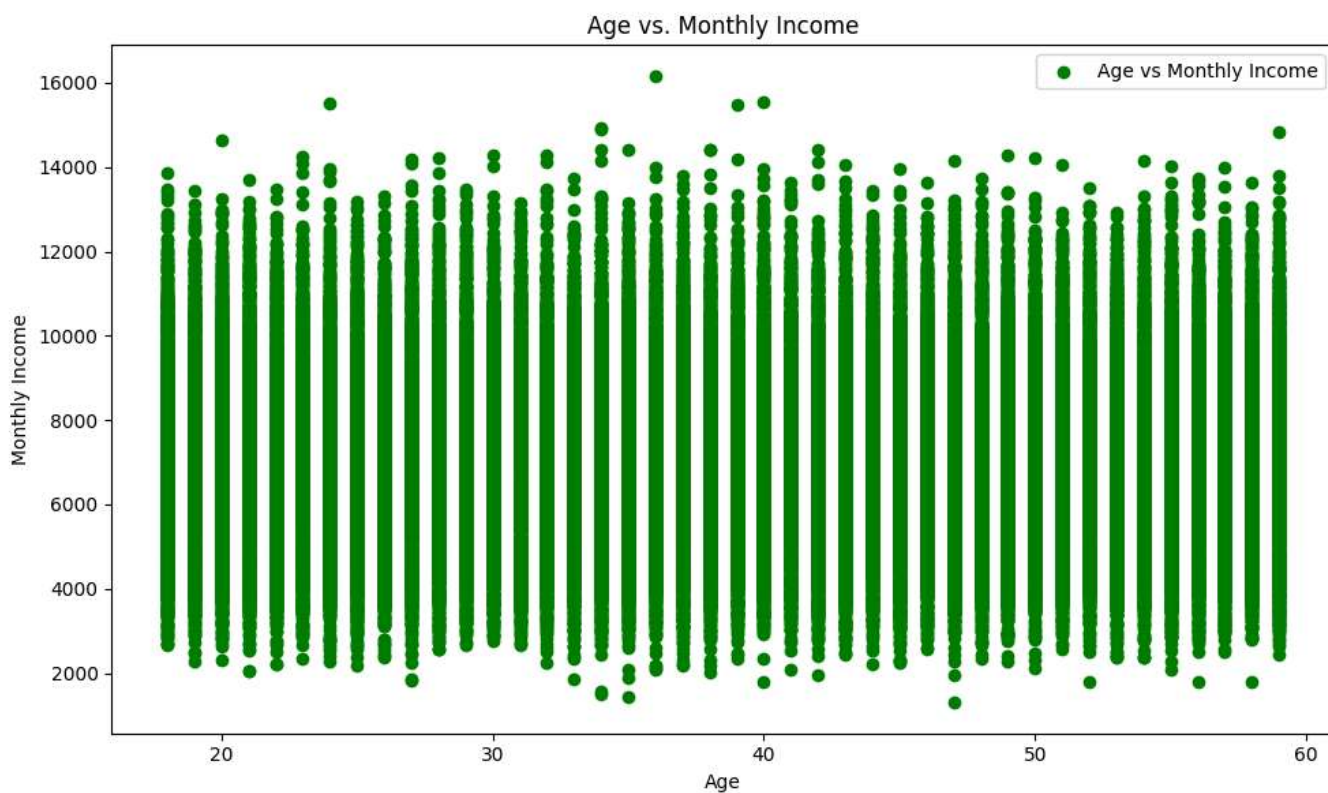
```
# 1. Line Plot - Example for Monthly Income over Years at Company
plt.figure(figsize=(10,6))
plt.plot(df['Years at Company'], df['Monthly Income'], marker='o', linestyle='-', color='blue')
plt.title('Monthly Income over Years at Company')
plt.xlabel('Years at Company')
plt.ylabel('Monthly Income')
```

```
plt.grid(True)
plt.tight_layout()
plt.show()
```



Employees with more years at the company tend to have higher salaries, but growth is not always consistent. Some employees with fewer years may earn more, possibly due to job roles, promotions, or performance

```
# 2. Scatter Plot - Age vs. Monthly Income
plt.figure(figsize=(10,6))
plt.scatter(df['Age'], df['Monthly Income'], c='green', label='Age vs Monthly Income')
plt.title('Age vs. Monthly Income')
plt.xlabel('Age')
plt.ylabel('Monthly Income')
plt.legend()
plt.tight_layout()
plt.show()
```



# 3. Bar Plot (Seaborn) - Job Role vs. Monthly Income

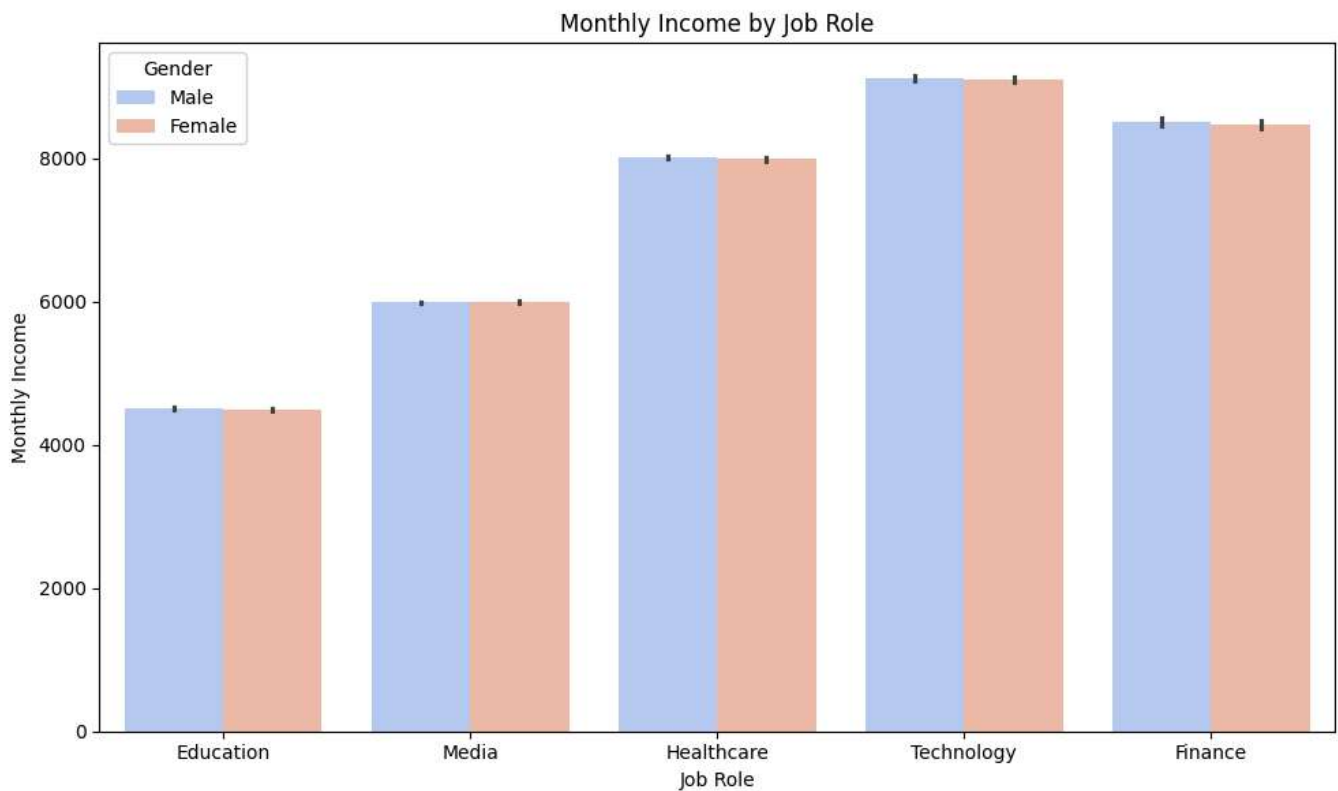
```
plt.figure(figsize=(10,6))
```

```
sns.barplot(x='Job Role', y='Monthly Income', data=df, palette='coolwarm',hue='Gender')
```

```
plt.title('Monthly Income by Job Role')
```

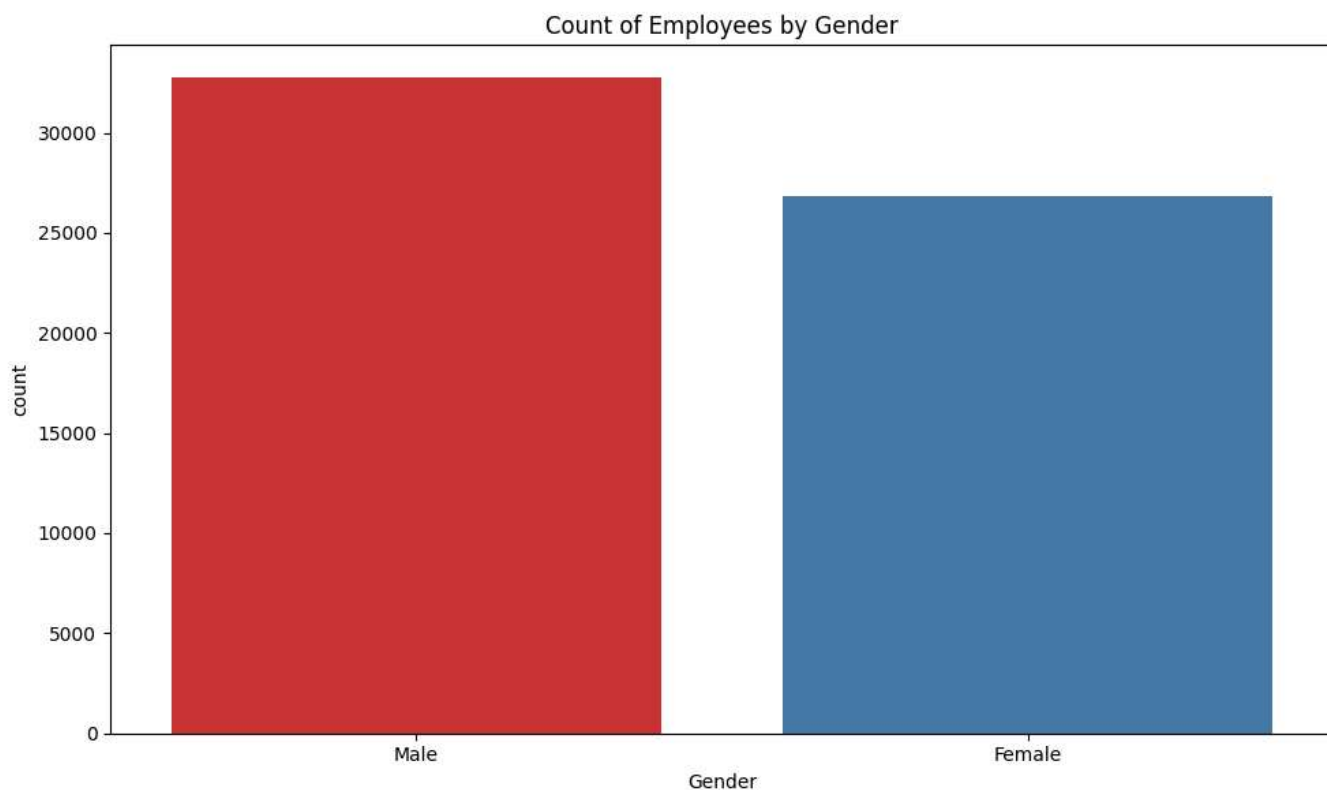
```
plt.tight_layout()
```

```
plt.show()
```



Certain job roles (e.g., Technology, Healthcare) likely have higher average salaries than others (e.g., Education, Media)

```
# 4. Count Plot (Seaborn) - Count of Employees by Gender
plt.figure(figsize=(10,6))
sns.countplot(x='Gender', data=df, palette='Set1', hue='Gender', legend=False)
plt.title('Count of Employees by Gender')
plt.tight_layout()
plt.show()
```



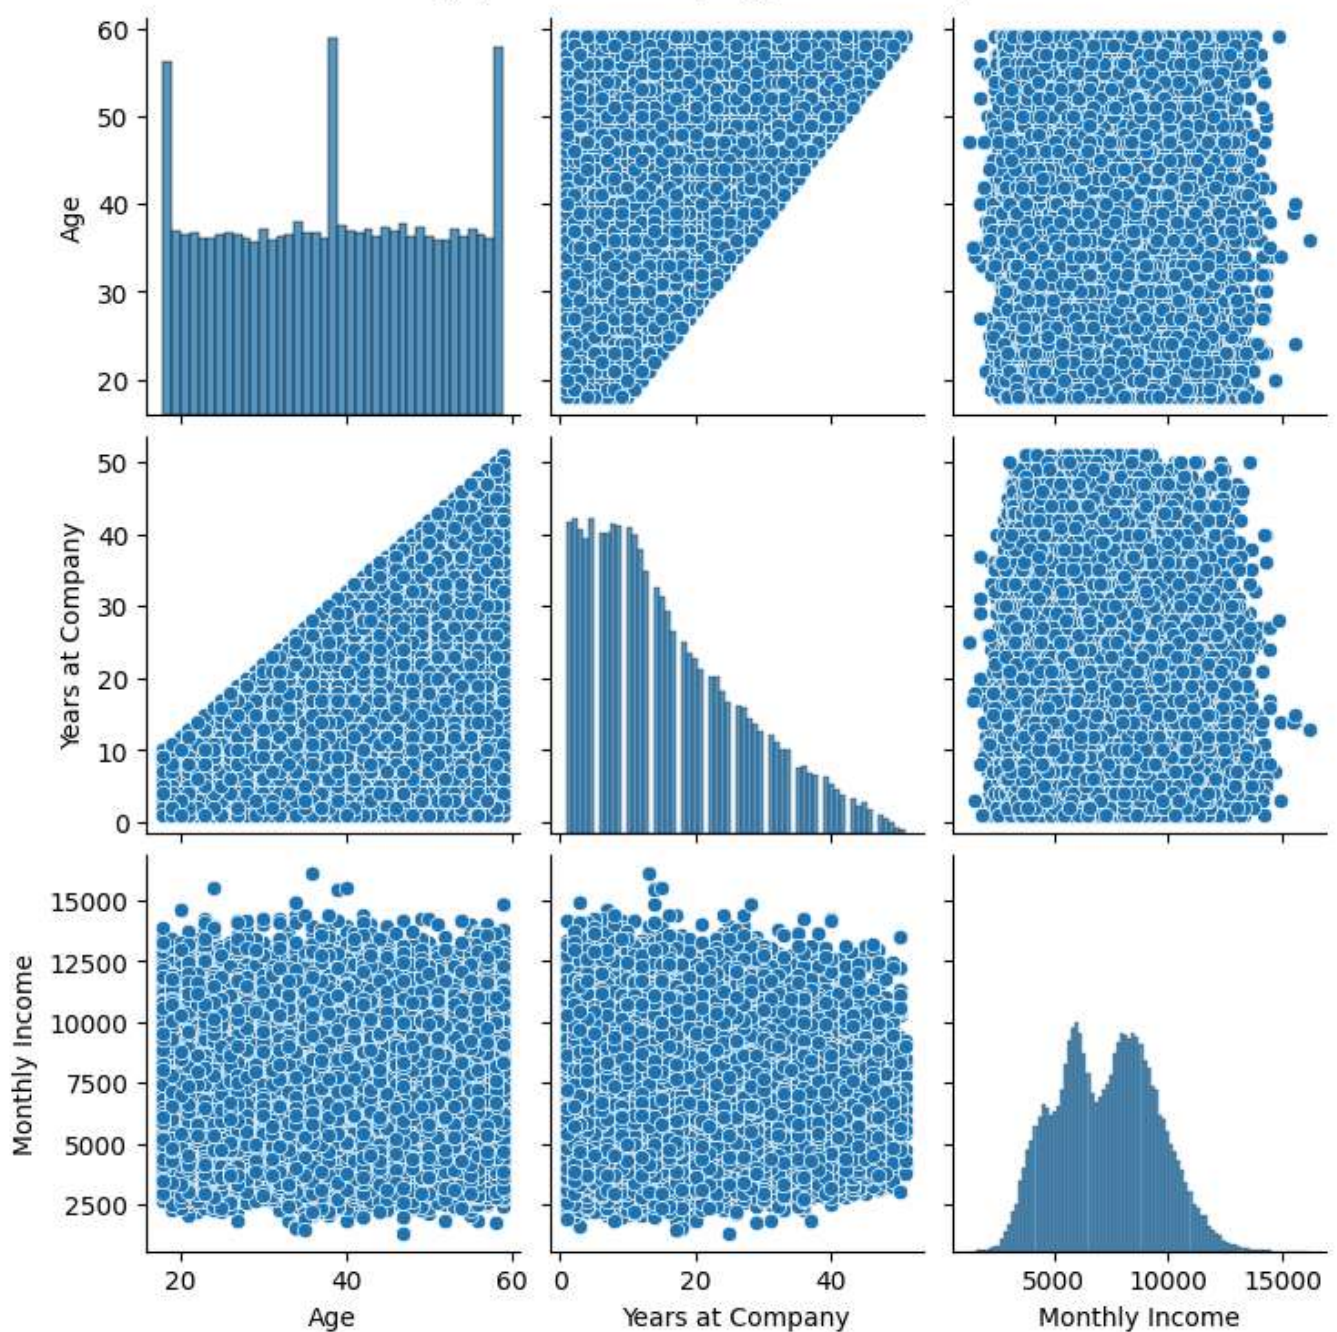
Shows the number of males and females present

```
# 5. Pair Plot - Relationship between multiple variables
sns.pairplot(df[['Age', 'Years at Company', 'Monthly Income']])
plt.suptitle('Pair Plot of Age, Years at Company, and Monthly Income', y=1.02)
plt.show()
```





Pair Plot of Age, Years at Company, and Monthly Income



#6.Shows if poor worklife balance leads to higher attrition

```
plt.figure(figsize=(8,5))
```


```
sns.countplot(x="Work-Life Balance", hue="Attrition", data=df, palette="Set2")
```

```
plt.title("Work-Life Balance vs. Attrition")
```

```
plt.xlabel("Work-Life Balance")
```

```
plt.ylabel("Count")
```

```
plt.show()
```



Drop Employee ID

```
df.drop(['Employee ID'],axis=1,inplace=True)
```

Label Encoding

```
features=df.select_dtypes(include=['object']).columns.tolist()
features
```




[Show hidden output](#)

```
encoder=LabelEncoder()
for col in features:
    df[col]=encoder.fit_transform(df[col])
df
```



	Age	Gender	Years at Company	Job Role	Monthly Income	Work- Life Balance	Job Satisfaction	Performance Rating	Number of Promotions
0	31	1	19	0	5390	0	2	0	
1	59	0	4	3	5534	3	0	3	
2	24	0	10	2	8159	2	0	3	
3	36	0	7	0	3989	2	0	2	
4	56	1	41	0	4821	1	3	0	
...	...	...	...	...	...	...	...	...	...
59593	50	0	12	0	4414	1	0	0	
59594	18	1	4	2	8040	1	0	2	
59595	22	0	14	4	7944	1	0	2	
59596	23	1	8	0	2931	1	3	0	
59597	56	1	19	4	6660	2	0	0	

59598 rows × 23 columns




```

# Compute the correlation matrix
correlation_matrix=df.corr()
# Create a heatmap to visualize the correlation matrix
plt.figure(figsize=(12, 8)) # Adjust figure size
sns.heatmap(correlation_matrix,annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
# Customize the plot
plt.title('Correlation_Heatmap', fontsize=16)
plt.show()

```





```
corr_matrix_unstacked = correlation_matrix.unstack() # Convert matrix to a Series
corr_matrix_unstacked = corr_matrix_unstacked[corr_matrix_unstacked != 1] # Remove self-cor
```

```
highest_corr = corr_matrix_unstacked.idxmax()
lowest_corr = corr_matrix_unstacked.idxmin()
```

```
print(f"Highest correlation: {highest_corr} -> {corr_matrix_unstacked.max():.4f}")
print(f"Lowest correlation: {lowest_corr} -> {corr_matrix_unstacked.min():.4f}")
```




```
Highest correlation: ('Job Role', 'Monthly Income') -> 0.5484
Lowest correlation: ('Marital Status', 'Attrition') -> -0.2296
```

```
X=df.iloc[:, :-1] #Input
X
```



	Age	Gender	Years at Company	Job Role	Monthly Income	Work- Life Balance	Job Satisfaction	Performance Rating	Number of Promotions
0	31	1	19	0	5390	0	2	0	
1	59	0	4	3	5534	3	0	3	
2	24	0	10	2	8159	2	0	3	
3	36	0	7	0	3989	2	0	2	
4	56	1	41	0	4821	1	3	0	
...	...	...	...	...	...	...	...	...	...
59593	50	0	12	0	4414	1	0	0	
59594	18	1	4	2	8040	1	0	2	
59595	22	0	14	4	7944	1	0	2	
59596	23	1	8	0	2931	1	3	0	
59597	56	1	19	4	6660	2	0	0	

59598 rows × 22 columns



```
y=df.iloc[:, -1]
y
```



### Attrition

<b>0</b>	1
<b>1</b>	1
<b>2</b>	1
<b>3</b>	1
<b>4</b>	1
...	...
<b>59593</b>	0
<b>59594</b>	0
<b>59595</b>	1
<b>59596</b>	0
<b>59597</b>	1

59598 rows × 1 columns

**dtype:** int64

## Scaling

```
scaler=MinMaxScaler()
X_scaled=scaler.fit_transform(X)
```

## Training and testing

```
X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=.2,random_state=42)
```

## Model Training

```
knn=KNeighborsClassifier()
svc=SVC()
gnb=GaussianNB()
dt=DecisionTreeClassifier()
rf=RandomForestClassifier()
gb=GradientBoostingClassifier()
ab=AdaBoostClassifier()
xg=XGBClassifier()
models=[knn,svc,gnb,dt,rf,gb,ab,xg]
for model in models:
    print('*****',model,'*****')
```

```

model.fit(X_train,y_train)
y_pred=model.predict(X_test)
print(classification_report(y_test,y_pred,digits=4))

```



```

***** KNeighborsClassifier() *****
precision    recall  f1-score   support

```

```

    0      0.6578    0.6524    0.6551     5667
    1      0.6873    0.6925    0.6899     6253

```

```

accuracy                0.6734     11920
macro avg      0.6726    0.6724    0.6725     11920
weighted avg   0.6733    0.6734    0.6733     11920

```

```

***** SVC() *****

```

```

precision    recall  f1-score   support

```

```

    0      0.7176    0.7081    0.7129     5667
    1      0.7386    0.7475    0.7430     6253

```

```

accuracy                0.7288     11920
macro avg      0.7281    0.7278    0.7279     11920
weighted avg   0.7286    0.7288    0.7287     11920

```

```

***** GaussianNB() *****

```

```

precision    recall  f1-score   support

```

```

    0      0.6642    0.7410    0.7005     5667
    1      0.7378    0.6605    0.6970     6253

```

```

accuracy                0.6987     11920
macro avg      0.7010    0.7007    0.6987     11920
weighted avg   0.7028    0.6987    0.6986     11920

```

```

***** DecisionTreeClassifier() *****

```

```

precision    recall  f1-score   support

```

```

    0      0.6373    0.6391    0.6382     5667
    1      0.6721    0.6704    0.6713     6253

```

```

accuracy                0.6555     11920
macro avg      0.6547    0.6548    0.6547     11920
weighted avg   0.6556    0.6555    0.6556     11920

```

```

***** RandomForestClassifier() *****

```

```

precision    recall  f1-score   support

```

```

    0      0.7243    0.7277    0.7260     5667
    1      0.7522    0.7489    0.7505     6253

```

```

accuracy                0.7388     11920
macro avg      0.7382    0.7383    0.7383     11920
weighted avg   0.7389    0.7388    0.7389     11920

```

```

***** GradientBoostingClassifier() *****

```

```

precision    recall  f1-score   support

```

0	0.7402	0.7392	0.7397	5667
1	0.7639	0.7649	0.7644	6253
accuracy			0.7527	11920
macro avg	0.7521	0.7521	0.7521	11920

```
selector = SelectKBest(score_func=f_classif, k=15) # Select Top 10 Features using ANOVA F-t
X_selected = selector.fit_transform(X,y)
selected_features = X.columns[selector.get_support()]
selected_features
```

```
Index(['Age', 'Gender', 'Years at Company', 'Work-Life Balance',
      'Job Satisfaction', 'Performance Rating', 'Number of Promotions',
      'Overtime', 'Distance from Home', 'Education Level', 'Marital Status',
      'Number of Dependents', 'Job Level', 'Remote Work',
      'Company Reputation'],
      dtype='object')
```

```
X_new=X[selected_features]
scaler=MinMaxScaler()
X_new_scaled=scaler.fit_transform(X_new)
```

```
y.value_counts()
```

```
count
Attrition
1      31260
0      28338

dtype: int64
```

## Sampling

```
from imblearn.over_sampling import SMOTE
os=SMOTE()
X_os,y_os=os.fit_resample(X_new_scaled,y)
```

```
y_os.value_counts()
```





**count**

**Attrition**

```
X_scaled=scaler.fit_transform(X_os)
```